

## ▼ Linear supervised regression

---

### 0. Import library

---

Import library

```
1 # Import libraries
2
3 # math library
4 import numpy as np
5
6 # visualization library
7 %matplotlib inline
8 from IPython.display import set_matplotlib_formats
9 set_matplotlib_formats('png2x','pdf')
10 import matplotlib.pyplot as plt
11
12 # machine learning library
13 from sklearn.linear_model import LinearRegression
14
15 # 3d visualization
16 from mpl_toolkits.mplot3d import axes3d
17
18 # computational time
19 import time
20
```

### ▼ 1. Load dataset

---

Load a set of data pairs  $\{x_i, y_i\}_{i=1}^n$  where  $x$  represents label and  $y$  represents target.

```
1 # import data with numpy
2 data = np.loadtxt('profit_population.txt', delimiter=',')
3
```


## ▼ 2. Explore the dataset distribution

---

Plot the training data points.

```
1 x_train = data[:,0]
2 y_train = data[:,1]
3
4 plt.title('Training data')
5 plt.xlabel('Population size (x 10$)')
6 plt.ylabel('Profit (x 10$)')
7 plt.scatter(x_train, y_train, c = 'r', marker = 'x')
```

↗ <matplotlib.collections.PathCollection at 0x7fab6a068208>



### ▼ 3. Define the linear prediction function

---

$$f_w(x) = w_0 + w_1 x$$

Vectorized implementation:

$$f_w(x) = Xw$$

with

$$X = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \\ 1 & x_n \end{bmatrix} \quad \text{and} \quad w = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} \quad \Rightarrow \quad f_w(x) = Xw = \begin{bmatrix} w_0 + w_1 x_1 \\ w_0 + w_1 x_2 \\ \vdots \\ w_0 + w_1 x_n \end{bmatrix}$$

```
1 # construct data matrix
2 X = np.ones((len(x_train), 2))
3 for i in range(len(X)):
4     X[i][1] = x_train[i]
5
6 # parameters vector
7 w = [1, 1]
8
9 # predictive function definition
10 def f_pred(X,w):
11
12     f = np.dot(X,w)
13
14     return f
15
16 # Test predictive function
17 y_pred = f_pred(X,w)
18
```

## ▼ 4. Define the linear regression loss

---

$$L(w) = \frac{1}{n} \sum_{i=1}^n \left( f_w(x_i) - y_i \right)^2$$

Vectorized implementation:

$$L(w) = \frac{1}{n} (Xw - y)^T (Xw - y)$$

with

$$Xw = \begin{bmatrix} w_0 + w_1 x_1 \\ w_0 + w_1 x_2 \\ \vdots \\ w_0 + w_1 x_n \end{bmatrix} \quad \text{and} \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

Implement the vectorized version of the linear regression loss function

```
1 # loss function definition
2 def loss_mse(y_pred,y):
3
4     loss = np.sum((y_pred - y) ** 2) / len(y)
5     return loss
6
7
8 # Test loss function
9 y = y_train
10 y_pred = f_pred(X,w)
11 loss = loss_mse(y_pred,y)
```

## ▼ 5. Define the gradient of the linear regression loss

---

Vectorized implementation: Given the loss

$$L(w) = \frac{1}{n}(Xw - y)^T(Xw - y)$$

The gradient is given by

$$\frac{\partial}{\partial w} L(w) = \frac{2}{n} X^T (Xw - y)$$

Implement the vectorized version of the gradient of the linear regression loss function.

```

1 # gradient function definition
2 def grad_loss(y_pred,y,X):
3
4     grad = np.dot(X.T, y_pred - y) / len(X) * 2
5
6     return grad
7
8
9 # Test grad function
10 y_pred = f_pred(X,w)
11 grad = grad_loss(y_pred,y,X)

```

## ▼ 6. Implement the gradient descent algorithm

---

- Vectorized implementation:

$$w^{k+1} = w^k - \tau \frac{2}{n} X^T (Xw^k - y)$$

Implement the vectorized version of the gradient descent function.

Plot the loss values  $L(w^k)$  with respect to iteration  $k$  the number of iterations.

```

1 # gradient descent function definition
2 def grad_desc(X, y, w_init, tau, max_iter):

```

```

3
4     L_iters = [] # record the loss values
5     w_iters = [] # record the parameter values
6     w = w_init # initialization
7
8     for i in range(max_iter): # loop over the iterations
9         tmp = []
10        y_pred = f_pred(X,w) # linear prediction function
11        grad_f = grad_loss(y_pred, y, X) # gradient of the loss
12
13        for j in range(len(w)):
14            w[j] = w[j] - tau * grad_f[j] # update rule of gradient descent
15            tmp.append(w[j])
16
17        w_iters.append(tmp) # save the current w value
18        L_iters.append(loss_mse(y_pred, y)) # save the current loss value
19
20    return w, L_iters, w_iters
21
22
23 # run gradient descent algorithm
24 start = time.time()
25 w_init = [0, 0]
26 tau = 0.01
27 max_iter = 50
28
29 w, L_iters, w_iters = grad_desc(X, y, w_init, tau, max_iter)
30
31 print('Time=',time.time() - start) # plot the computational cost
32 print(L_iters[-1] ) # plot the last value of the loss
33 print(w) # plot the last value of the parameter w
34
35
36 # plot
37 plt.figure(2)
38 plt.plot([i for i in range(len(L_iters))], L_iters) # plot the loss curve
39 plt.xlabel('Iterations')
40 plt.ylabel('Loss')

```

```


10 plt.figure(figsize=(10, 5))
41 plt.show()

```

```

Time= 0.0014889240264892578
10.973835031833165
[-0.5770974029560515, 0.8596358917014932]

```




## ▼ 7. Plot the linear prediction function

$$f_w(x) = w_0 + w_1 x$$

```

1 # linear regression model
2 x_pred = np.linspace(0,25,100) # define the domain of the prediction function
3 y_pred = x_pred * w[1] + w[0]
4 # plot
5 plt.figure(3)
6 plt.scatter(x_train, y_train, c='r', marker = 'x')
7 plt.plot(x_pred, y_pred, label = 'gradient descent optimization' )
8 plt.legend(loc='best')
9 plt.title('Training data')
10 plt.xlabel('Population size (x 10k)')
11 plt.ylabel('Profit $(x 10k)')
12 plt.show()

```



## ▼ 8. Comparison with Scikit-learn linear regression algorithm

### Compare with the Scikit-learn solution

```

1 # run linear regression with scikit-learn
2 start = time.time()
3 lin_reg_sklearn = LinearRegression()
4 sklearn_x_train = x_train.reshape(len(x_train), 1)
5 sklearn_y_train = y_train.reshape(len(y_train), 1)
6
7 lin_reg_sklearn.fit(sklearn_x_train,sklearn_y_train) # learn the model parameters
8 print('Time=',time.time() - start)
9
10
11 # compute loss value
12 w_sklearn = np.zeros([2,1])
13 w_sklearn[0,0] = lin_reg_sklearn.intercept_
14 w_sklearn[1,0] = lin_reg_sklearn.coef_
15

```



```

15
16 print(w_sklearn)
17
18 loss_sklearn = lin_reg_sklearn.score(sklearn_x_train, sklearn_y_train) # compute the loss from the sklearn solution
19
20 print('loss sklearn=', loss_sklearn)
21 print('loss gradient descent=', L_iters[-1])
22
23
24 # plot
25 y_pred_sklearn = lin_reg_sklearn.predict(np.linspace(0,25,100).reshape(100, 1))

```

```

↳ Time= 0.001645803451538086
   [[-3.89578088]
    [ 1.19303364]]
   loss sklearn= 0.7020315537841397
   loss gradient descent= 10.973835031833165

```

```


1 plt.figure(3)
2
3 plt.scatter(x_train, y_train, c='r', marker = 'x')
4 plt.plot(x_pred, y_pred, label = 'gradient descent optimization' )
5 plt.plot(x_pred, y_pred_sklearn, c = 'g' ,label = 'Scikit-learn optimization' )
6 plt.legend(loc='best')
7 plt.title('Training data')
8 plt.xlabel('Population size (x 10k)')
9 plt.ylabel('Profit $(x 10k)')
10 plt.show()

```

```

↳

```



## ▼ 9. Plot the loss surface, the contours of the loss and the gradient descent steps

```



1 # Create grid coordinates for plotting a range of  $L(w_0, w_1)$ -values
2 B0 = np.linspace(-10, 10, 50)
3 B1 = np.linspace(-1, 4, 50)
4
5 xx, yy = np.meshgrid(B0, B1, indexing='xy')
6 Z = np.zeros((B0.size, B1.size))
7
8 # Calculate loss values based on  $L(w_0, w_1)$ -values
9 for (i, j), v in np.ndenumerate(Z):
10     contour_pred = x_train * B1[i] + B0[j]
11     Z[i, j] = np.sum((contour_pred - y) ** 2) / len(y)
12 # 3D visualization
13 fig = plt.figure(figsize=(10, 3))
14 ax1 = fig.add_subplot(121)
15 ax2 = fig.add_subplot(122, projection='3d')
16
17 # Left plot
18 CS = ax1.contour(xx, yy, Z, np.logspace(-2, 3, 20), cmap=plt.cm.jet)
19 ax1.scatter(w[0], w[1])
20 ax1.plot(np.array(w_iters)[:, 0], np.array(w_iters)[:, 1]) # plot the loss curve
21
22 # Right plot
23 ax2.plot_surface(xx, yy, Z, rstride=1, cstride=1, alpha=0.6, cmap=plt.cm.jet)
24 ax2.set_zlabel('Loss  $L(w_0, w_1)$ ')
25 ax2.set_zlim(Z.min(), Z.max())
26 ax2.plot(np.array(w_iters)[:, 0], np.array(w_iters)[:, 1])

```

```

27 ax2.scatter(w[0], w[1])
28
29 # settings common to both plots
30 for ax in fig.axes:
31     ax.set_xlabel(r'$w_0$', fontsize=17)
32     ax.set_ylabel(r'$w_1$', fontsize=17)

```



## ▼ Output results

### ▼ 1. Plot the training data (1pt)


```

1 plt.title('Training data')
2 plt.xlabel('Population size (x 10$)')
3 plt.ylabel('Profit (x 10$)')
4 plt.scatter(x_train, y_train, c = 'r', marker = 'x')

```




<matplotlib.collections.PathCollection at 0x7fab654a5668>



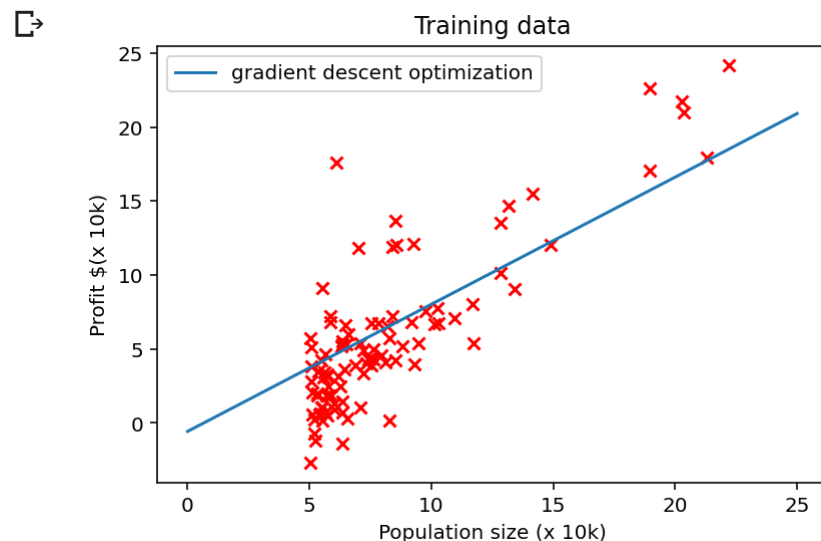
## ▼ 2. Plot the loss curve in the course of gradient descent (2pt)

```
1 plt.figure(2)
2 plt.plot([i for i in range(len(L_iters))], L_iters) # plot the loss curve
3 plt.xlabel('Iterations')
4 plt.ylabel('Loss')
5 plt.show()
```



### ▼ 3. Plot the prediction function superimposed on the training data (2pt)

```
1 plt.scatter(x_train, y_train, c='r', marker = 'x')
2 plt.plot(x_pred, y_pred, label = 'gradient descent optimization' )
3 plt.legend(loc='best')
4 plt.title('Training data')
5 plt.xlabel('Population size (x 10k)')
6 plt.ylabel('Profit $(x 10k)')
7 plt.show()
```




### ▼ 4. Plot the prediction functions obtained by both the Scikit-learn linear regression solution and the gradient descent superimposed on the training data (2pt)

```
1 plt.scatter(x_train, y_train, c='r', marker = 'x')
2 plt.plot(x_pred, y_pred, label = 'gradient descent optimization' )
3 plt.plot(x_pred, y_pred_sklearn, c = 'g', label = 'Scikit-learn optimization' )
```

```

4 plt.legend(loc='best')
5 plt.title('Training data')
6 plt.xlabel('Population size (x 10k)')
7 plt.ylabel('Profit $(x 10k)')
8 plt.show()

```




## ▼ 5. Plot the loss surface (right) and the path of the gradient descent (2pt)

```

1 fig = plt.figure(figsize=(15,6))
2 ax2 = fig.add_subplot(122, projection='3d')
3 ax2.plot_surface(xx, yy, Z, rstride=1, cstride=1, alpha=0.6, cmap=plt.cm.jet)
4 ax2.set_zlabel('Loss  $L(w_0, w_1)$ ')
5 ax2.set_zlim(Z.min(), Z.max())
6 ax2.plot(np.array(w_iters)[: ,0], np.array(w_iters)[: ,1])
7 ax2.scatter(w[0], w[1])
8
9 for ax in fig.axes:
10     ax.set_xlabel(r'$w_0$', fontsize=17)
11     ax.set_ylabel(r'$w_1$', fontsize=17)

```



▼ 6. Plot the contour of the loss surface (left) and the path of the gradient descent (2pt)

```
1 fig = plt.figure(figsize=(15,6))
2 ax1 = fig.add_subplot(121)
3
4 CS = ax1.contour(xx, yy, Z, np.logspace(-2, 3, 20), cmap=plt.cm.jet)
5 ax1.scatter(w[0], w[1])
6 ax1.plot(np.array(w_iters)[: ,0], np.array(w_iters)[: ,1]) # plot the loss curve
7
8 for ax in fig.axes:
9     ax.set_xlabel(r'$w_0$', fontsize=17)
10    ax.set_ylabel(r'$w_1$', fontsize=17)
```



